

```
# programmer - Sophia Quinton
# date - 11-3-21
# class - DSC -540
# assignment - Assignment 1
```

#libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sns
from sklearn import datasets, linear_model
from sklearn.model_selection import train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn import metrics
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
```

#Part1 - Tools Readiness

```
##file from (Larose & Larose, 2019)
```

```
##pandas
```

```
frame = pd.read_csv("E:/GCU/Graduate Classes/DSC - 540 Machine Learning/Week 1/cereals.csv")
```

```
frame.head()
```

	Name	Manuf	Type	Calories	Protein	Fat	Sodium	\
0	100%_Bran	N	C	70	4	1	130	
1	100%_Natural_Bran	Q	C	120	3	5	15	
2	All-Bran	K	C	70	4	1	260	
3	All-Bran_with_Extra_Fiber	K	C	50	4	0	140	
4	Almond_Delight	R	C	110	2	2	200	

	Fiber	Carbo	Sugars	...	Weight	Cups	Rating	Cold	Nabisco	Quaker
0	10.0	5.0	6.0	...	1.0	0.33	68.402973	1	1	0
1	2.0	8.0	8.0	...	1.0	1.00	33.983679	1	0	1
2	9.0	7.0	5.0	...	1.0	0.33	59.425505	1	0	0
3	14.0	8.0	0.0	...	1.0	0.50	93.704912	1	0	0
4	1.0	14.0	8.0	...	1.0	0.75	34.384843	1	0	0

	Kelloggs	GeneralMills	Ralston	AHFP
0	0	0	0	0
1	0	0	0	0
2	1	0	0	0
3	1	0	0	0
4	0	0	1	0

```
[5 rows x 23 columns]
```

```
##numpy
```


```
rounded_rating = np.round(frame['Rating'][0], 3)  
rounded_rating
```

```
68.403
```

```
##matplotlib
```

```
plt.boxplot(frame['Rating'])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x20c3812e5e0>,  
             <matplotlib.lines.Line2D at 0x20c3812e8b0>],  
 'caps': [<matplotlib.lines.Line2D at 0x20c3812ec40>,  
          <matplotlib.lines.Line2D at 0x20c3812efd0>],  
 'boxes': [<matplotlib.lines.Line2D at 0x20c3812e190>],  
 'medians': [<matplotlib.lines.Line2D at 0x20c3813e3a0>],  
 'fliers': [<matplotlib.lines.Line2D at 0x20c3813e730>],  
 'means': []}
```



```
##scikit-learn
```

```
frame_train, frame_test = train_test_split(frame, test_size=0.2,  
                                             random_state=25)  
frame_train.head()  
print(len(frame_train))  
print(len(frame_test))
```

```
61
```

```
16
```

```
#Part2 - Review Predictive Models and Python Proficiency
```

```
##read in data
```

```
estate = pd.read_csv("E:/GCU/Graduate Classes/DSC - 540 Machine Learning/Week
```

```
1/housing.csv")
estate.tail()
```


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
\					
20635	-121.09	39.48	25.0	1665.0	374.0
20636	-121.21	39.49	18.0	697.0	150.0
20637	-121.22	39.43	17.0	2254.0	485.0
20638	-121.32	39.43	18.0	1860.0	409.0
20639	-121.24	39.37	16.0	2785.0	616.0

	population	households	median_income	median_house_value	\
20635	845.0	330.0	1.5603	78100.0	
20636	356.0	114.0	2.5568	77100.0	
20637	1007.0	433.0	1.7000	92300.0	
20638	741.0	349.0	1.8672	84700.0	
20639	1387.0	530.0	2.3886	89400.0	


	ocean_proximity
20635	INLAND
20636	INLAND
20637	INLAND
20638	INLAND
20639	INLAND

```
## view data
```


```
plt.scatter(estate['longitude'], estate['latitude'])
plt.show()
```




```
plt.boxplot(estate['housing_median_age'])  
plt.show()
```




```
plt.boxplot(estate['total_rooms'])  
plt.show()
```




```
plt.boxplot(estate['population'])  
plt.show()
```



```
plt.boxplot(estate['median_income'])  
plt.show()
```



```
plt.boxplot(estate['median_house_value'])  
plt.show()
```



```
##split data
```

```
estate_train, estate_test = train_test_split(estate, test_size=0.2,
random_state=25)
```

```
##prepare data (GeeksforGeeks, 2020)
```

```
test_columns = estate[['longitude', 'latitude', 'housing_median_age',
'total_rooms', 'population', 'median_income']]
vif = pd.DataFrame()
vif["variable"] = test_columns.columns
vif["VIF"] = [variance_inflation_factor(test_columns.values,i) for i in
range(test_columns.shape[1])]
vif
```

	variable	VIF
0	longitude	592.503040
1	latitude	538.400848
2	housing_median_age	7.238474
3	total_rooms	11.505213
4	population	11.408182
5	median_income	6.264782

```
##do PCA Analysis
```

```
x_estate_train = estate_train[['longitude', 'latitude', 'housing_median_age',
'total_rooms', 'population', 'median_income']]
x_estate_test = estate_test[['longitude', 'latitude', 'housing_median_age',
'total_rooms', 'population', 'median_income']]
y_estate_train = estate_train['median_house_value']
y_estate_test = estate_test['median_house_value']
```

```
scaler = StandardScaler()
```

```

scaler.fit(x_estate_train)
x_estate_train = scaler.transform(x_estate_train)
x_estate_test = scaler.transform(x_estate_test)

pca = PCA(0.95)
pca.fit(x_estate_train)
x_estate_train = pca.transform(x_estate_train)
x_estate_test = pca.transform(x_estate_test)

##run model (Larose & Larose, 2019)
constantX = sm.add_constant(x_estate_train, prepend=True)
estate_model = sm.OLS(y_estate_train, constantX).fit()
estate_model.summary()

```

```

<class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
=
Dep. Variable:          median_house_value    R-squared:
0.528
Model:                                OLS    Adj. R-squared:
0.528
Method:                    Least Squares    F-statistic:
4619.
Date:                Mon, 01 Nov 2021    Prob (F-statistic):
0.00
Time:                20:11:37    Log-Likelihood:            -
2.0975e+05
No. Observations:                16512    AIC:
4.195e+05
Df Residuals:                16507    BIC:
4.196e+05
Df Model:                                4
Covariance Type:                nonrobust
=====
=
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-
const      2.066e+05    619.105     333.762     0.000    2.05e+05
2.08e+05
x1         7677.6593    418.833      18.331     0.000    6856.701
8498.617
x2          934.7505    458.486       2.039     0.041     36.068
1833.433
x3        -7.378e+04    615.501    -119.877     0.000    -7.5e+04    -
7.26e+04
x4         4.268e+04    695.343      61.375     0.000    4.13e+04

```

4.4e+04

```
=====
=
Omnibus:                3482.398    Durbin-Watson:
1.993
Prob(Omnibus):           0.000    Jarque-Bera (JB):
8769.875
Skew:                    1.158    Prob(JB):
0.00
Kurtosis:                5.718    Cond. No.
1.66
=====
=
```

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

##(DataScience+, nd)


```
map= pd.DataFrame(pca.components_,columns=['longitude', 'latitude',
'housing_median_age', 'total_rooms', 'population', 'median_income'])
plt.figure(figsize=(12,6))
sns.heatmap(map,cmap='twilight')
```

<AxesSubplot:>



##scree plot (Zach, 2021)

```
PC_values = np.arange(pca.n_components_) + 1
plt.plot(PC_values, pca.explained_variance_ratio_)
plt.show()
```

##validate data (Larose & Larose, 2019)

```
constantX_test = sm.add_constant(x_estate_test, prepend=True)
estate_model_test= sm.OLS(y_estate_test, constantX_test).fit()
estate_model_test.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

```
=====
```

```
=
```

```
Dep. Variable:    median_house_value    R-squared:
0.527
Model:                OLS    Adj. R-squared:
0.527
Method:             Least Squares    F-statistic:
1149.
Date:                Mon, 01 Nov 2021    Prob (F-statistic):
0.00
Time:                17:06:16    Log-Likelihood:
52369.
No. Observations:    4128    AIC:
1.047e+05
Df Residuals:        4123    BIC:
1.048e+05
Df Model:              4
Covariance Type:      nonrobust
```

```
=====
```

```
=
```

```
coef    std err        t    P>|t|    [0.025
0.975]
```

```

-----
-
const      2.092e+05   1218.673   171.633   0.000   2.07e+05
2.12e+05
x1         7460.1512   790.853    9.433   0.000   5909.652
9010.650
x2        -205.8746    893.178   -0.230   0.818  -1956.985
1545.235
x3       -7.419e+04   1231.663  -60.240   0.000  -7.66e+04  -
7.18e+04
x4        4.184e+04   1351.925   30.947   0.000   3.92e+04
4.45e+04
=====
=
Omnibus:                809.222   Durbin-Watson:
2.045
Prob(Omnibus):          0.000   Jarque-Bera (JB):
1785.669
Skew:                   1.123   Prob(JB):
0.00
Kurtosis:               5.310   Cond. No.
1.72
=====
=

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

```

ypred = estate_model.predict(constantX_test)
##MAE (Larose & Larose, 2019)
print("Model MAE: ", metrics.mean_absolute_error(y_true = y_estate_test,
y_pred = ypred))
print("base MAE: ", np.sqrt(metrics.mean_squared_error(y_true =
y_estate_test, y_pred = ypred)))

```

```

Model MAE:  58203.1309674163
base MAE:  78282.93963949986

```

```

# calculate s or standard deviation (Larose & Larose, 2019)
print("test MSE", np.sqrt(estate_model_test.scale))
print("Model MSE", np.sqrt(estate_model.scale))

```

```

test MSE 78268.50127486532
Model MSE 79554.41649151935

```